

*Operating Systems:  
Internals and Design Principles, 6/E*  
William Stallings

# Memory Management

Dr. Sanam Shahla Rizvi



# Roadmap

- ▶ Basic requirements of Memory Management
- ▶ Memory Partitioning
- ▶ Basic blocks of memory management
  - Paging
  - Segmentation

# Basic requirements of Memory Management



# The need for memory management

- ▶ Memory is cheap today, and getting cheaper
  - But applications are demanding more and more memory, there is never enough!
- ▶ Memory Management, involves swapping blocks of data from secondary storage.
- ▶ Memory I/O is slow compared to a CPU
  - The OS must cleverly time the swapping to maximise the CPU's efficiency

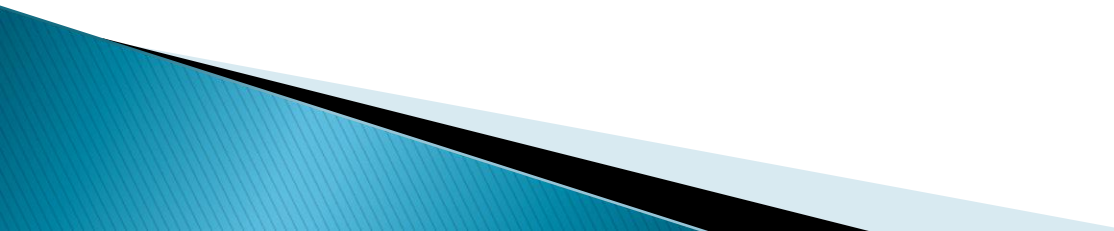
# Memory Management

*Memory needs to be allocated to ensure a reasonable supply of ready processes to consume available processor time*

# Memory Management Terms

Term	Description
Frame	<b><i>Fixed</i></b> -length block of main memory.
Page	<b><i>Fixed</i></b> -length block of data in secondary memory (e.g. on disk).
Segment	<b><i>Variable-length</i></b> block of data that resides in secondary memory.

# Memory Management Requirements

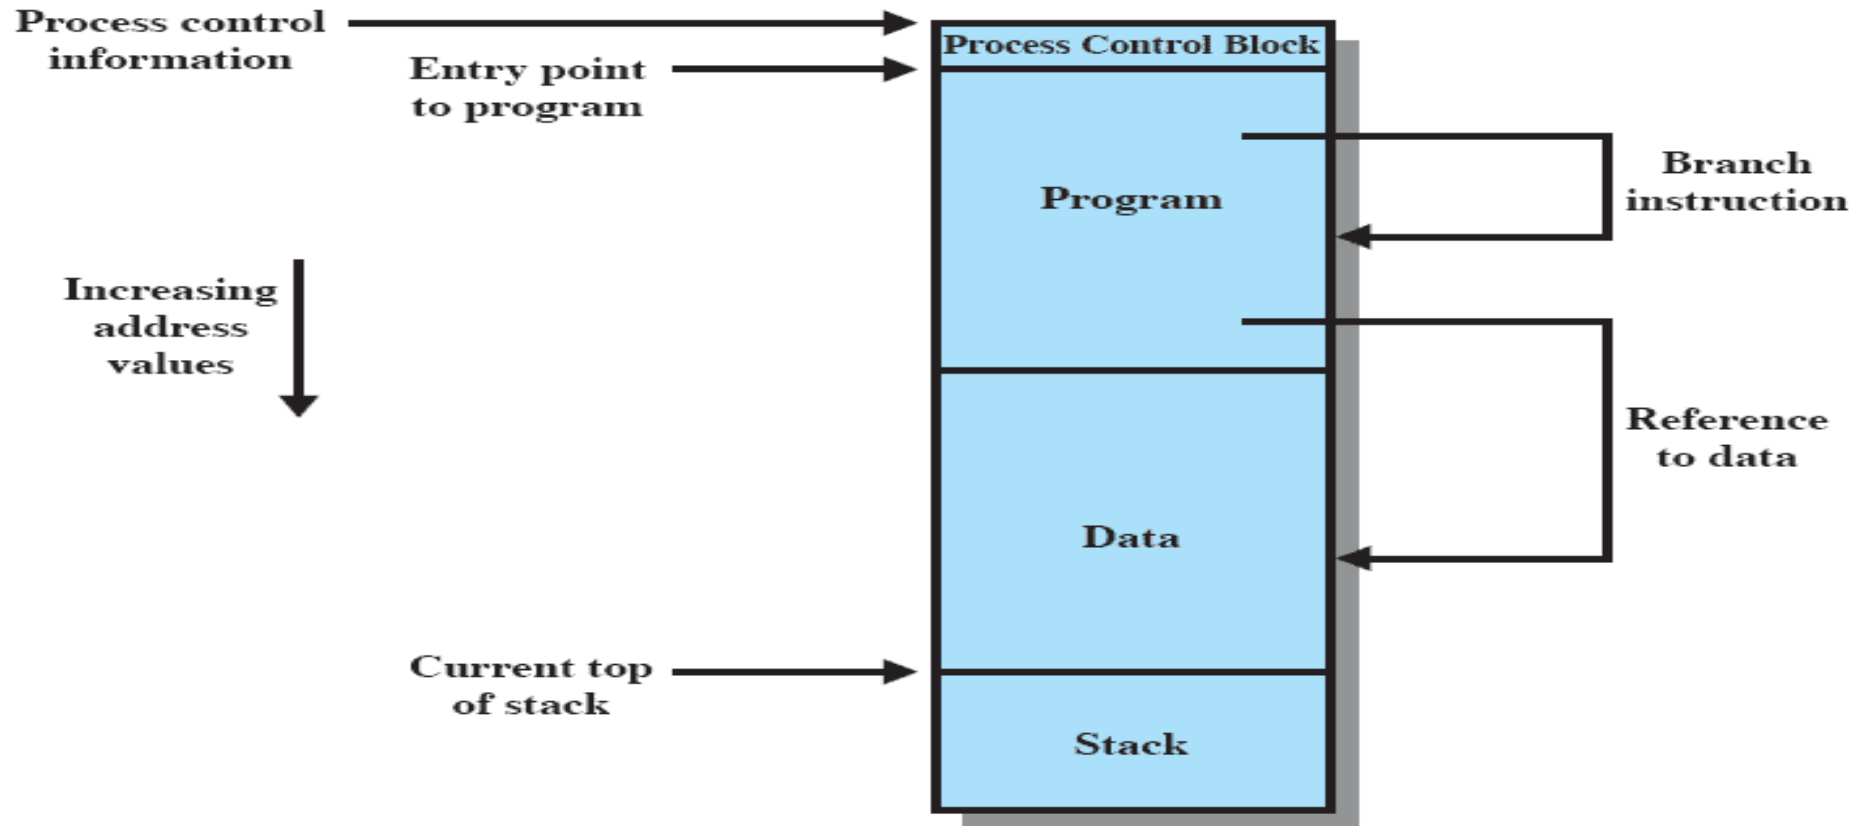
- ▶ Relocation
  - ▶ Protection
  - ▶ Sharing
  - ▶ Logical organisation
  - ▶ Physical organisation
- 

# Requirements: Relocation

- ▶ The programmer does not know where the program will be placed in memory when it is executed,
  - it may be swapped to disk and return to main memory at a different location (relocated)
- ▶ Memory references must be translated to the actual physical memory address

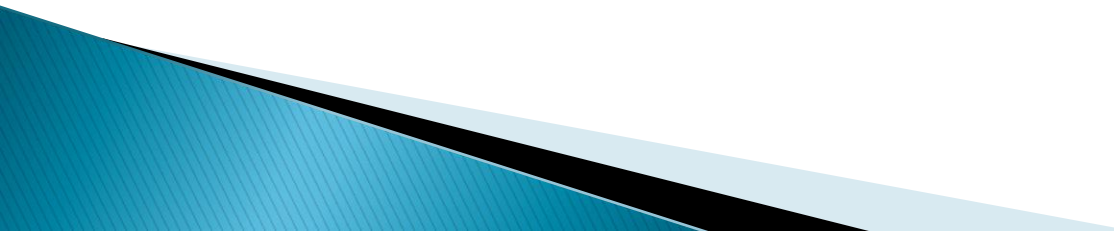


# Addressing



**Figure 7.1** Addressing Requirements for a Process

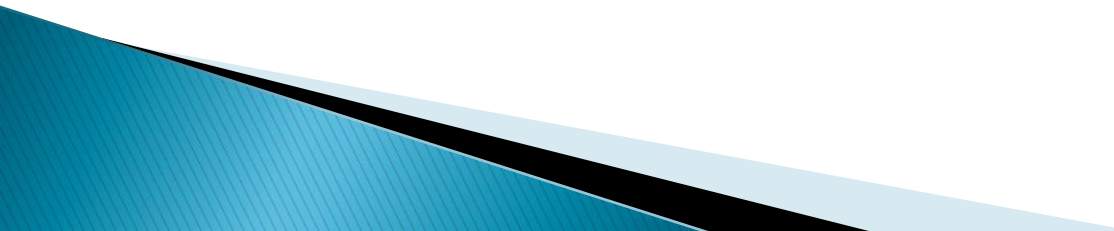
# Requirements: Protection

- ▶ Processes should not be able to reference memory locations in another process without permission
  - ▶ Operating system cannot anticipate all of the memory references that a program will make
  - ▶ Assess the permissibility of a memory reference (data access or branch) at the time of execution of the instruction making the reference
- 

# Requirements: Sharing

- ▶ Allow several processes to access the same portion of memory
- ▶ Better to allow each process access to the same copy of the program rather than have their own separate copy

# Requirements: Logical Organization

- ▶ Memory is organized linearly (usually)
  - ▶ Programs are written in modules
    - Modules can be written and compiled independently
  - ▶ Different degrees of protection given to modules (read-only, execute-only)
  - ▶ Share modules among processes
  - ▶ Segmentation helps here
- 

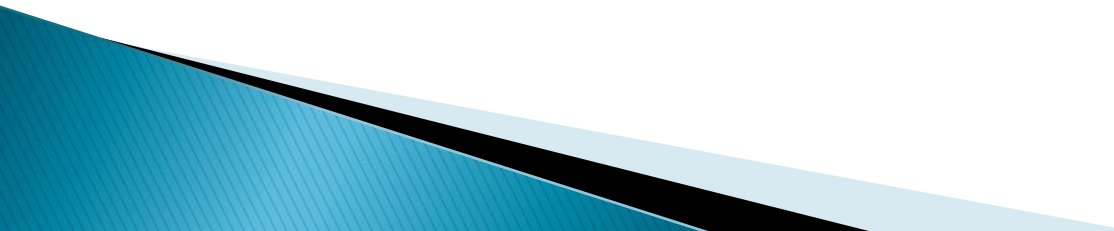
# Requirements: Physical Organization

- ▶ Cannot leave the programmer with the responsibility to manage memory
- ▶ Memory available for a program plus its data may be insufficient
  - Overlaying allows various modules to be assigned the same region of memory but is time consuming
- ▶ Programmer does not know how much space will be available

# Memory Partitioning




# Types of Partitioning

- ▶ Fixed Partitioning
  - ▶ Dynamic Partitioning
  - ▶ Simple Paging
  - ▶ Simple Segmentation
  - ▶ Virtual Memory Paging
  - ▶ Virtual Memory Segmentation
- 

# Fixed Partitioning

- ▶ Equal-size partitions
  - Any process whose size is less than or equal to the partition size can be loaded into an available partition
- ▶ The operating system can swap a process out of a partition
  - If none are in a ready or running state



(a) Equal-size partitions




# Fixed Partitioning Problems

- ▶ A program may not fit in a partition.
  - The programmer must design the program with overlays
- ▶ Main memory use is inefficient.
  - Any program, no matter how small, occupies an entire partition.
  - This results in *internal fragmentation*.

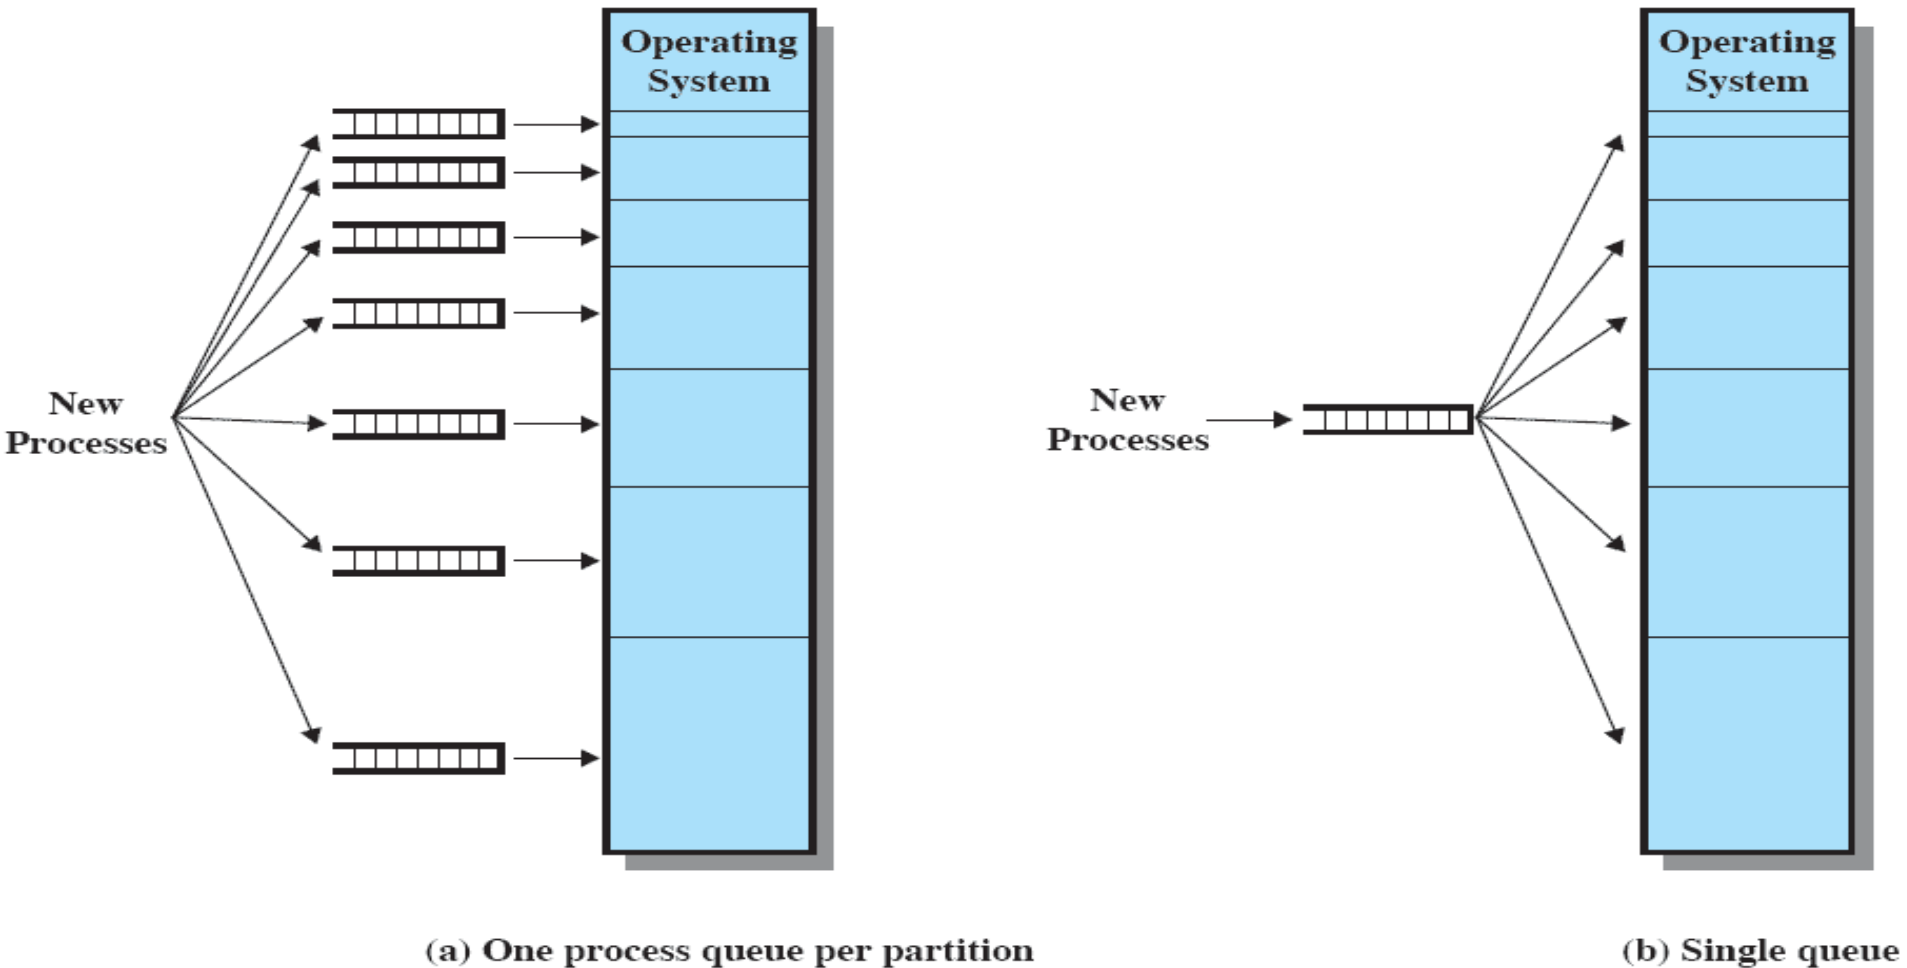
# Solution – Unequal Size Partitions

- ▶ Lessens both problems
  - but doesn't solve completely
  - Programs up to 16M can be accommodated without overlay
  - Smaller programs can be placed in smaller partitions, reducing internal fragmentation



(b) Unequal-size partitions

# Fixed Partitioning – Unequal Size



**Figure 7.3** Memory Assignment for Fixed Partitioning

# Dynamic Partitioning

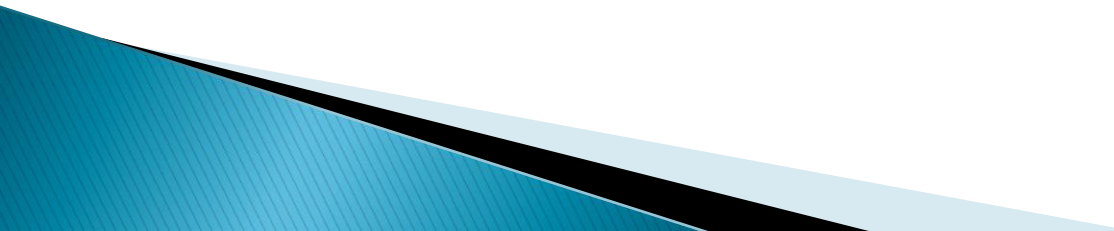
- ▶ Partitions are of variable length and number
- ▶ Process is allocated exactly as much memory as required

# Dynamic Partitioning Example



- ▶ ***External Fragmentation***
- ▶ Memory external to all processes is fragmented
- ▶ Can resolve using ***compaction***
  - OS moves processes so that they are contiguous
  - Time consuming and wastes CPU time

# Dynamic Partitioning

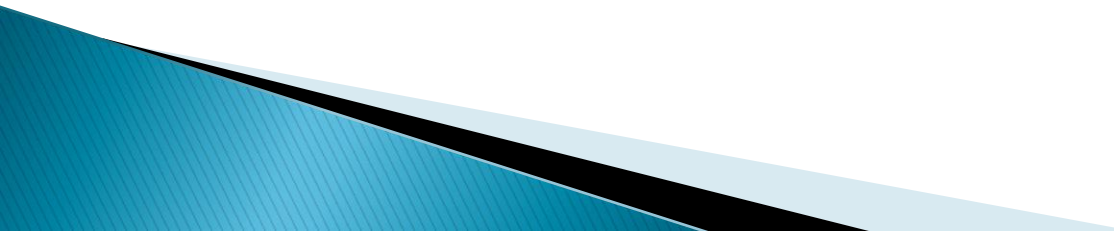
- ▶ Operating system must decide which free block to allocate to a process
  - ▶ Best-fit algorithm
    - Chooses the block that is closest in size to the request
    - Worst performance overall
    - Since smallest block is found for process, the smallest amount of fragmentation is left
    - Memory compaction must be done more often
- 

# Dynamic Partitioning

- ▶ First-fit algorithm
  - Scans memory from the beginning and chooses the first available block that is large enough
  - Fastest
  - May have many process loaded in the front end of memory that must be searched over when trying to find a free block

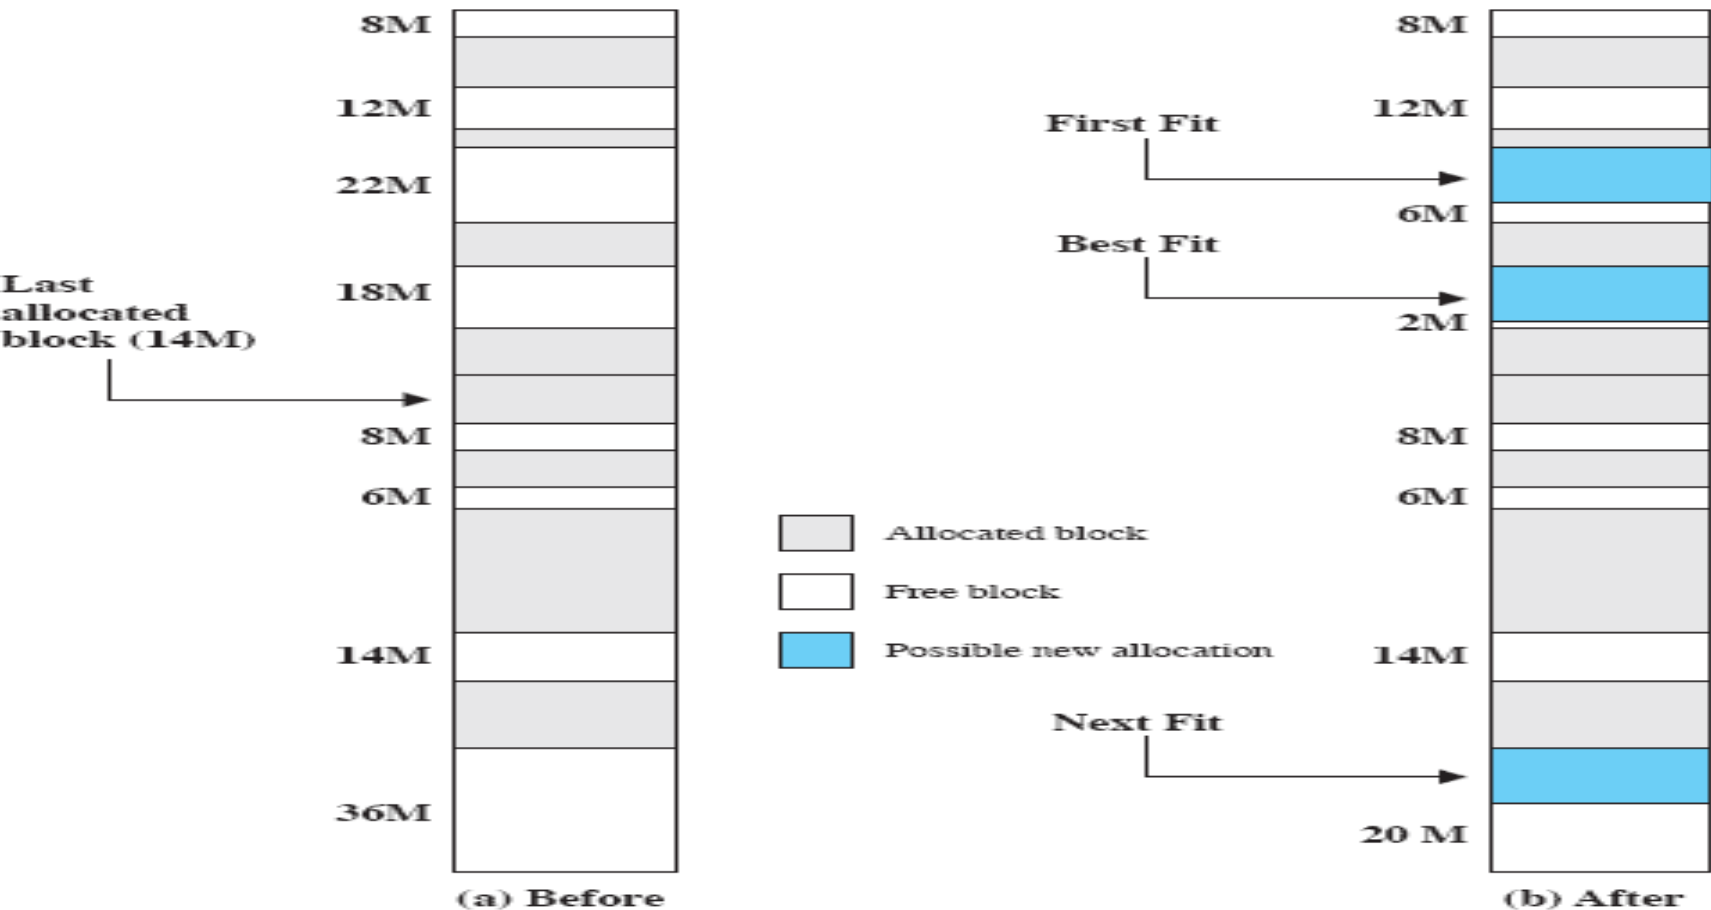
# Dynamic Partitioning

## ▶ Next-fit

- Scans memory from the location of the last placement
  - More often allocate a block of memory at the end of memory where the largest block is found
  - The largest block of memory is broken up into smaller blocks
  - Compaction is required to obtain a large block at the end of memory
- 



# Allocation



**Figure 7.5 Example Memory Configuration before and after Allocation of 16-Mbyte Block**

# Buddy System

- ▶ Entire space available is treated as a single block
- ▶ Request arrives of less than block size
- ▶ Block splits into two equal size blocks (buddies)
  - Process continues until smallest block greater than or equal to the request is generated

# Example of Buddy System

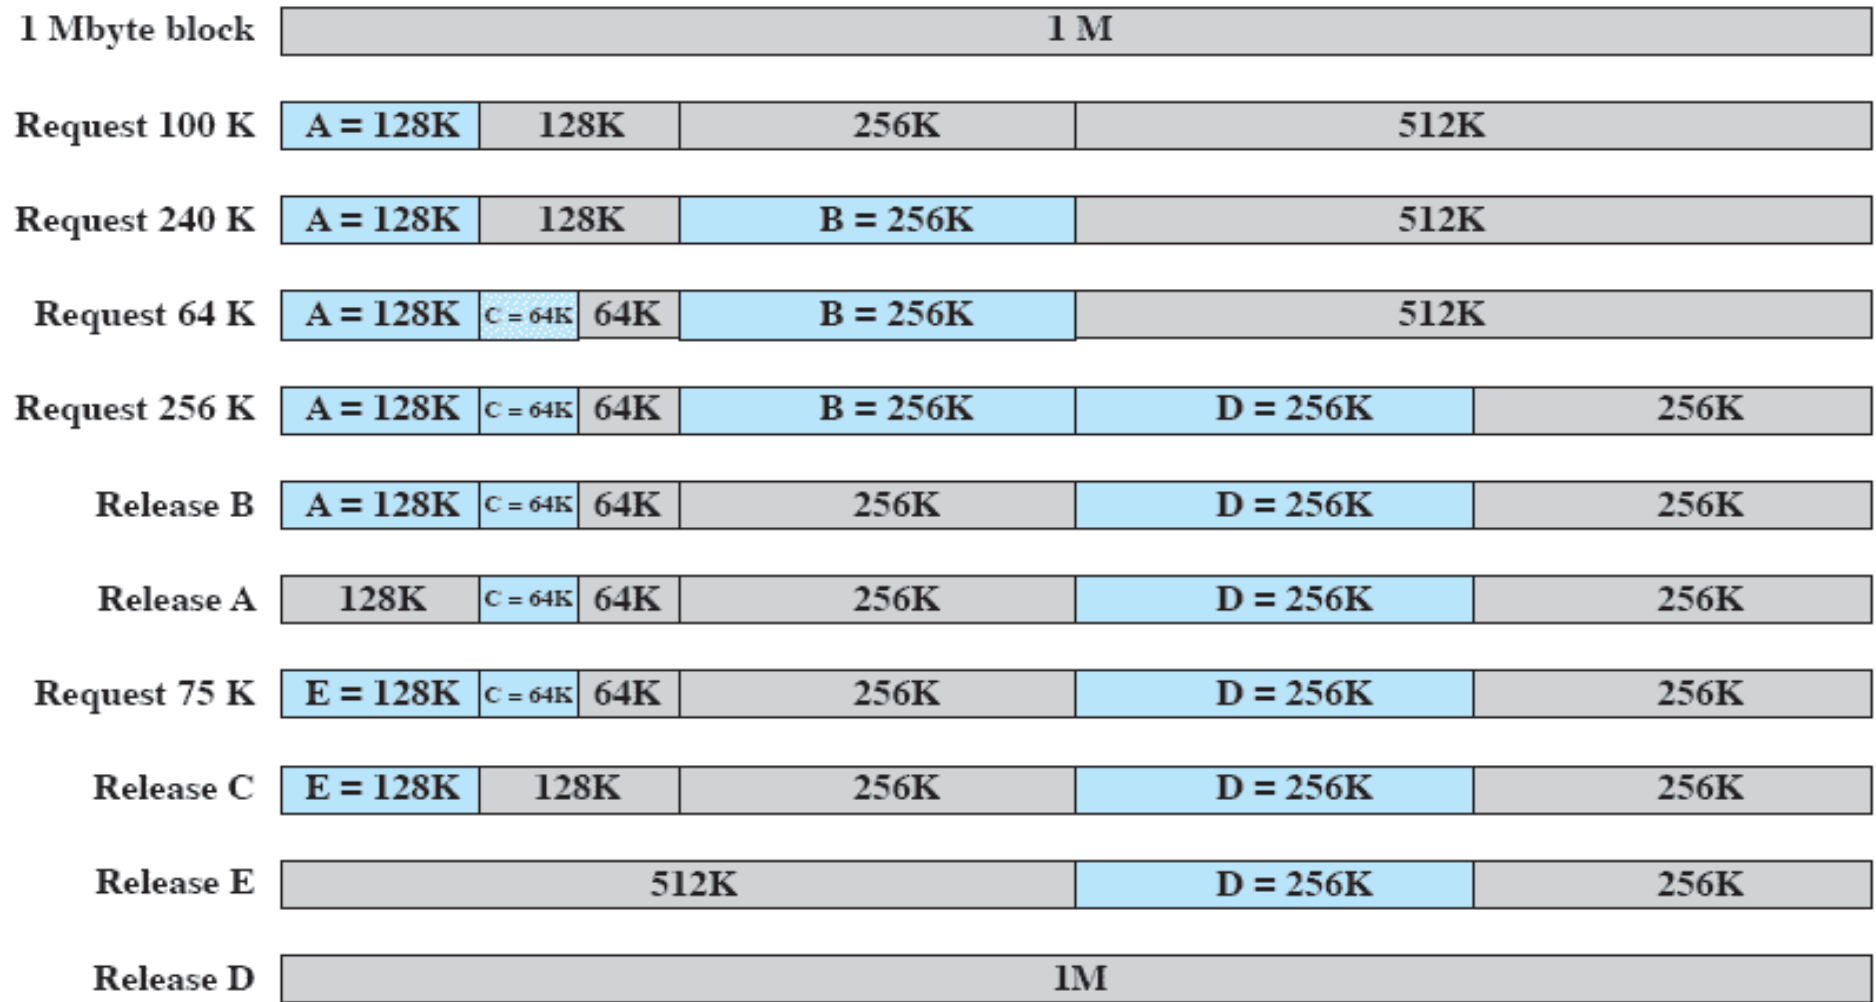



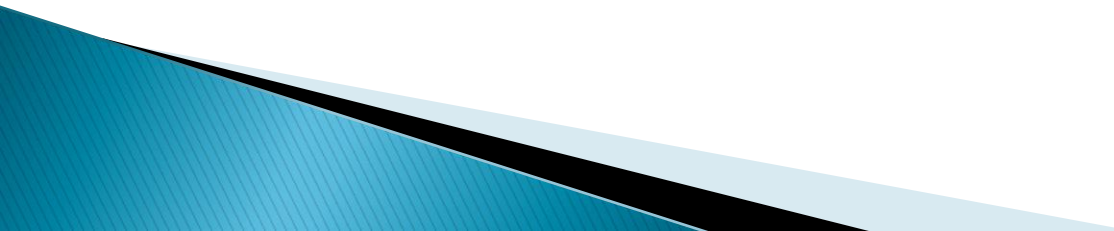
Figure 7.6 Example of Buddy System

# Tree Representation of Buddy System

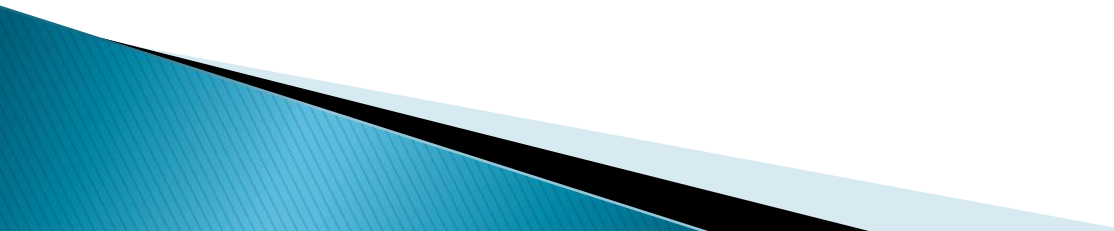


**Figure 7.7** Tree Representation of Buddy System

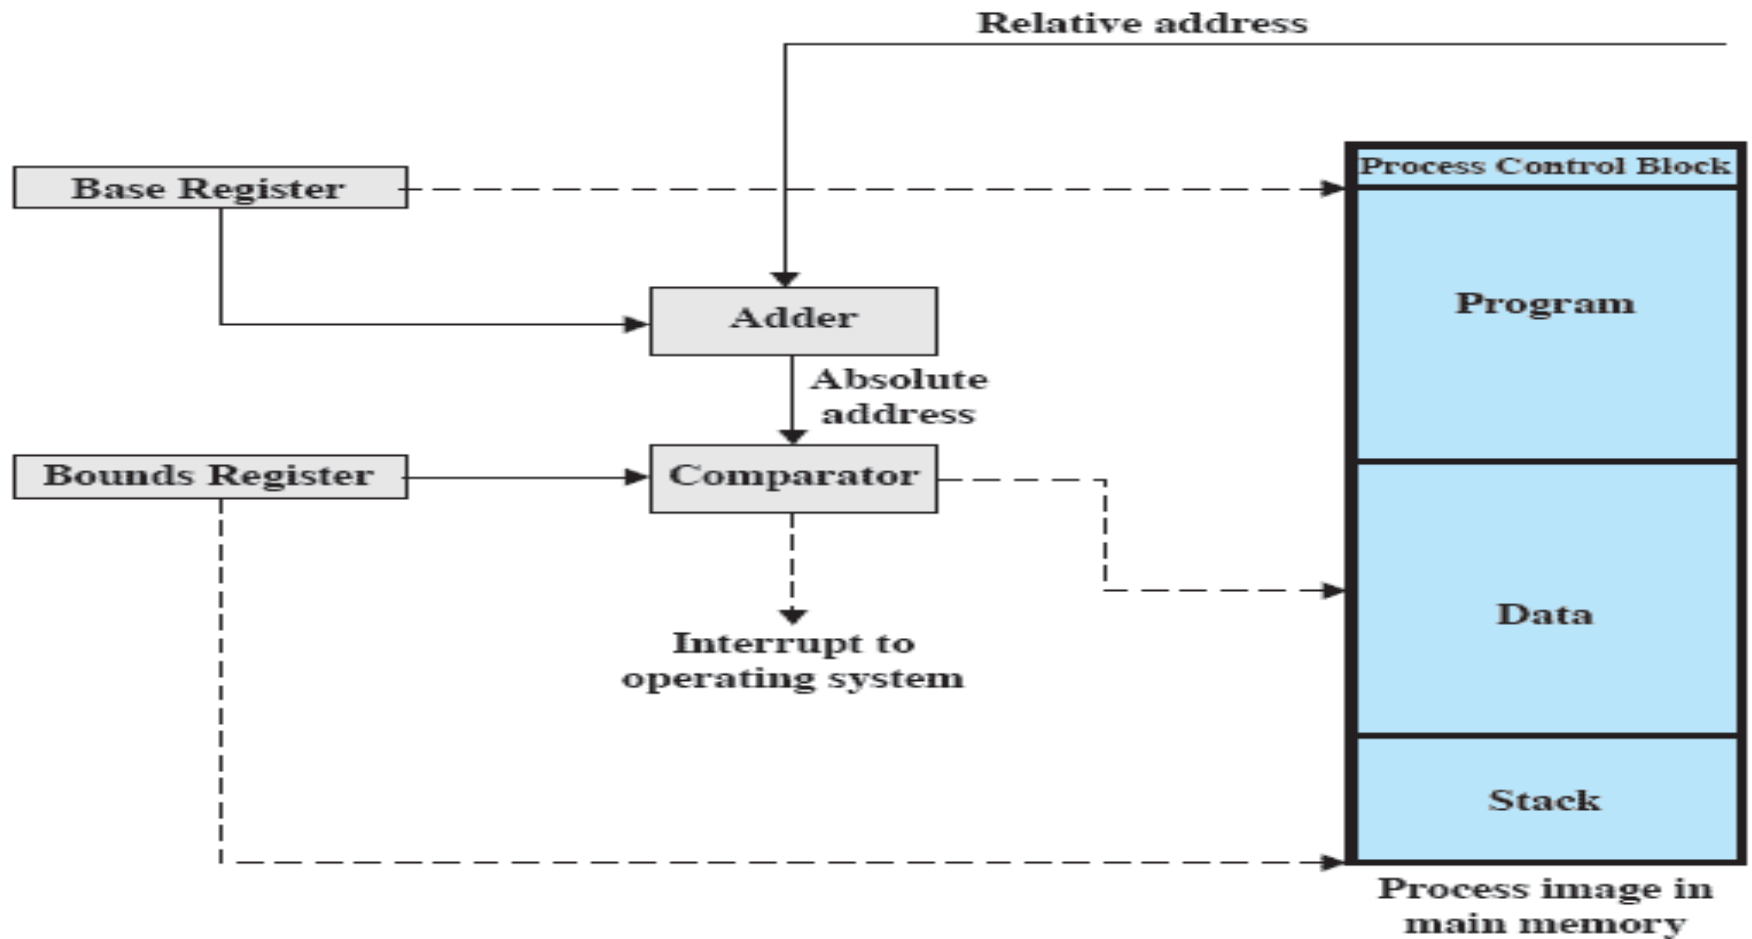
# Relocation

- ▶ When program loaded into memory the actual (absolute) memory locations are determined
  - ▶ A process may occupy different partitions which means different absolute memory locations during execution
    - Swapping
    - Compaction
- 

# Addresses

- ▶ Logical
    - Address is a reference to a memory location independent of the current assignment of data to memory.
  - ▶ Relative
    - Address is example of logical address.
    - Address expressed as a location relative to some known point usually a processor register.
  - ▶ Physical or Absolute
    - Address is actual location in main memory.
- 

# Relocation



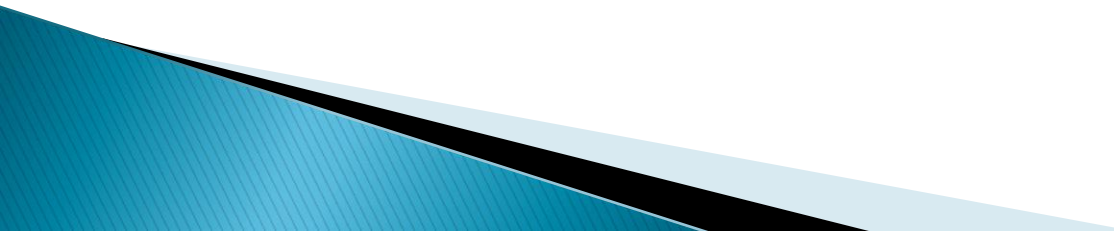
**Figure 7.8 Hardware Support for Relocation**

# Registers Used during Execution

- ▶ Base register
  - Starting address for the process
- ▶ Bounds register
  - Ending location of the process
- ▶ These values are set when the process is loaded or when the process is swapped in




# Registers Used during Execution

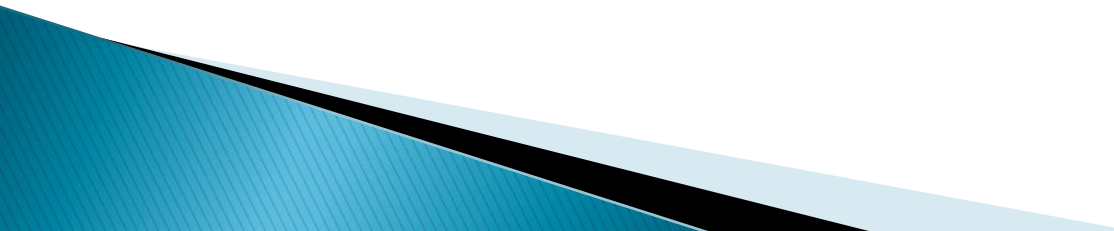
- ▶ The value of the base register is added to a relative address to produce an absolute address
  - ▶ The resulting address is compared with the value in the bounds register
  - ▶ If the address is not within bounds, an interrupt is generated to the operating system
- 

# Basic blocks of memory management

Paging  
Segmentation



# Paging

- ▶ Partition memory into small equal fixed-size chunks and divide each process into the same size chunks
  - ▶ The chunks of a process are called *pages*
  - ▶ The chunks of memory are called *frames*
- 

# Paging

- ▶ Operating system maintains a page table for each process
  - Contains the frame location for each page in the process
  - Memory address consist of a page number and offset within the page

# Page Table

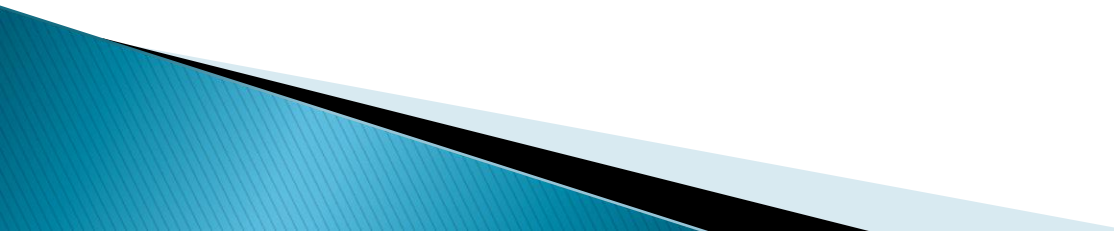
# Processes and Frames

0	0	0	-	0	7	0	4	13
1	1	1	-	1	8	1	5	14
2	2	2	-	2	9	2	6	
3	3			3	10	3	11	
						4	12	
Process A page table		Process B page table		Process C page table		Process D page table		Free frame list

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	D.0
5	D.1
6	D.2
7	C.0
8	C.1
9	C.2
10	C.3
11	D.3
12	D.4
13	
14	

Figure 7.10 Data Structures for the Example of Figure 7.9 at Time Epoch (f)

# Segmentation

- ▶ A program can be subdivided into segments
    - Segments may vary in length
    - There is a maximum segment length
  - ▶ Addressing consist of two parts
    - a segment number and
    - an offset
  - ▶ Segmentation is similar to dynamic partitioning
- 

# Logical Addresses


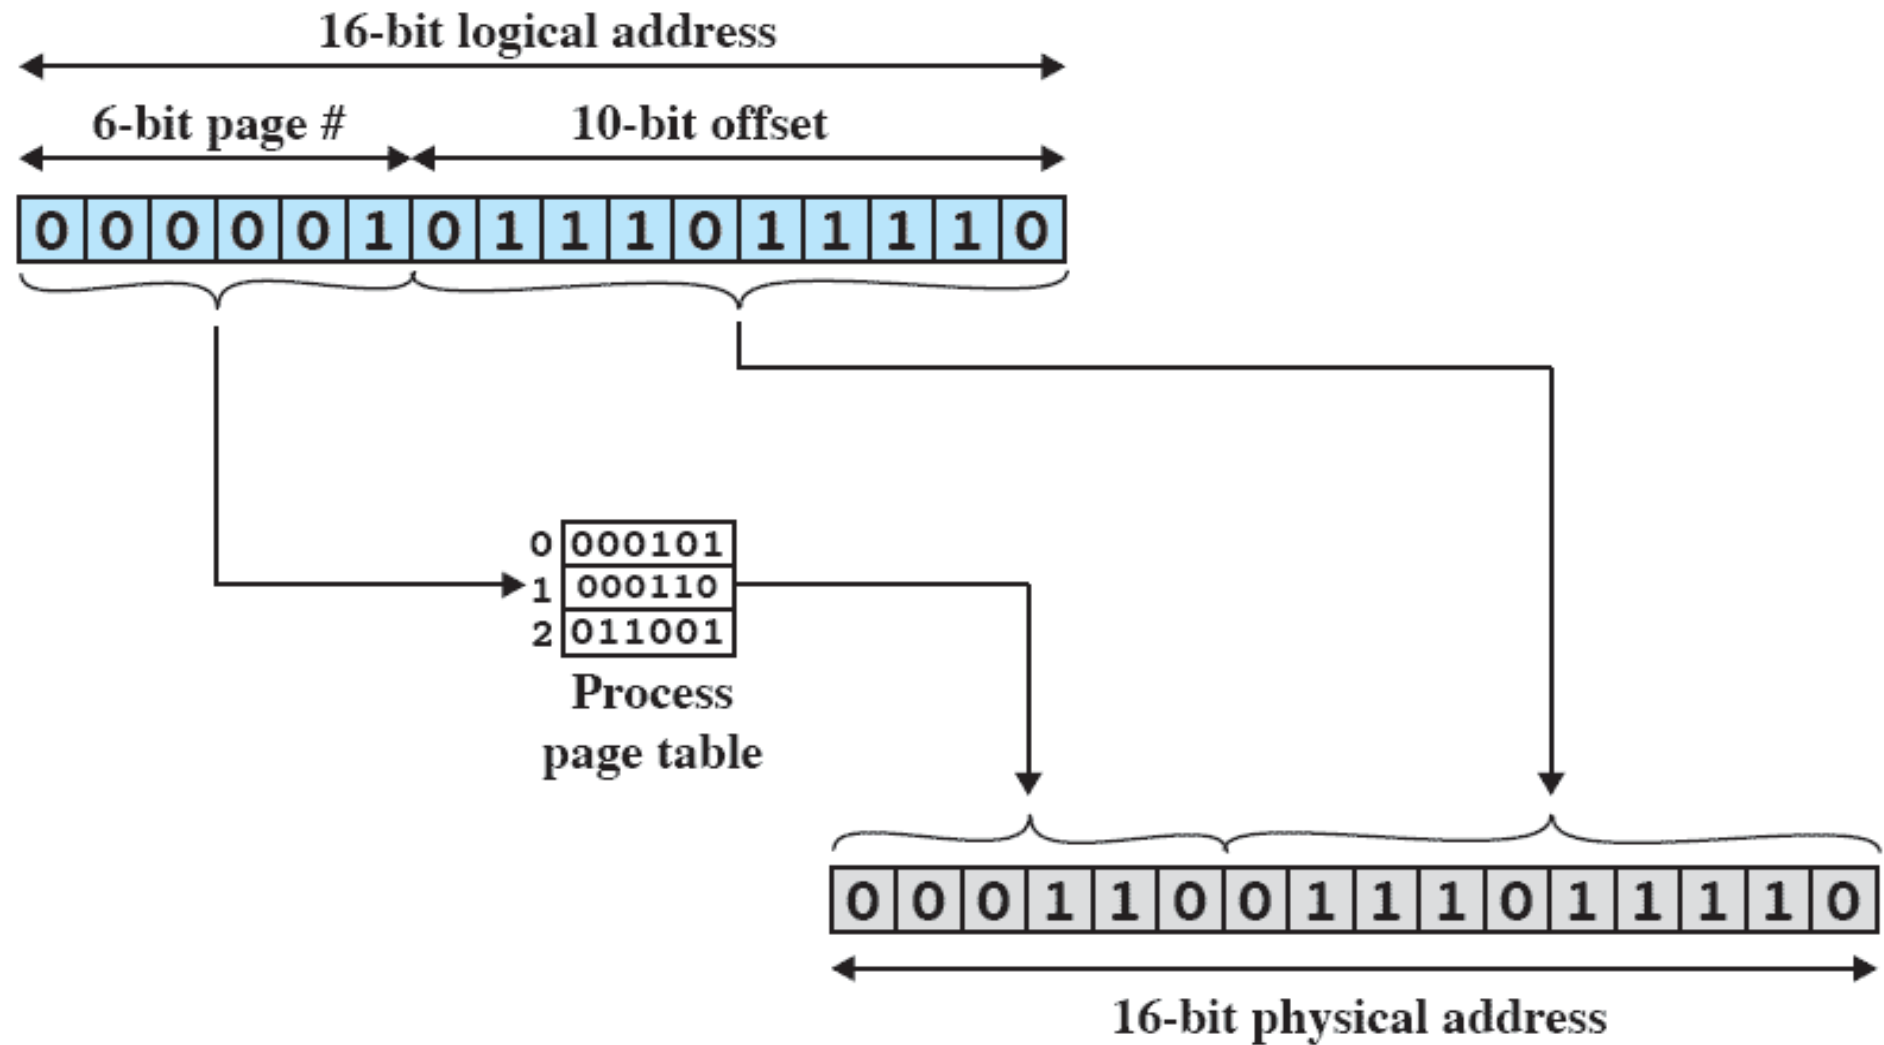


Figure 7.11 Logical Addresses

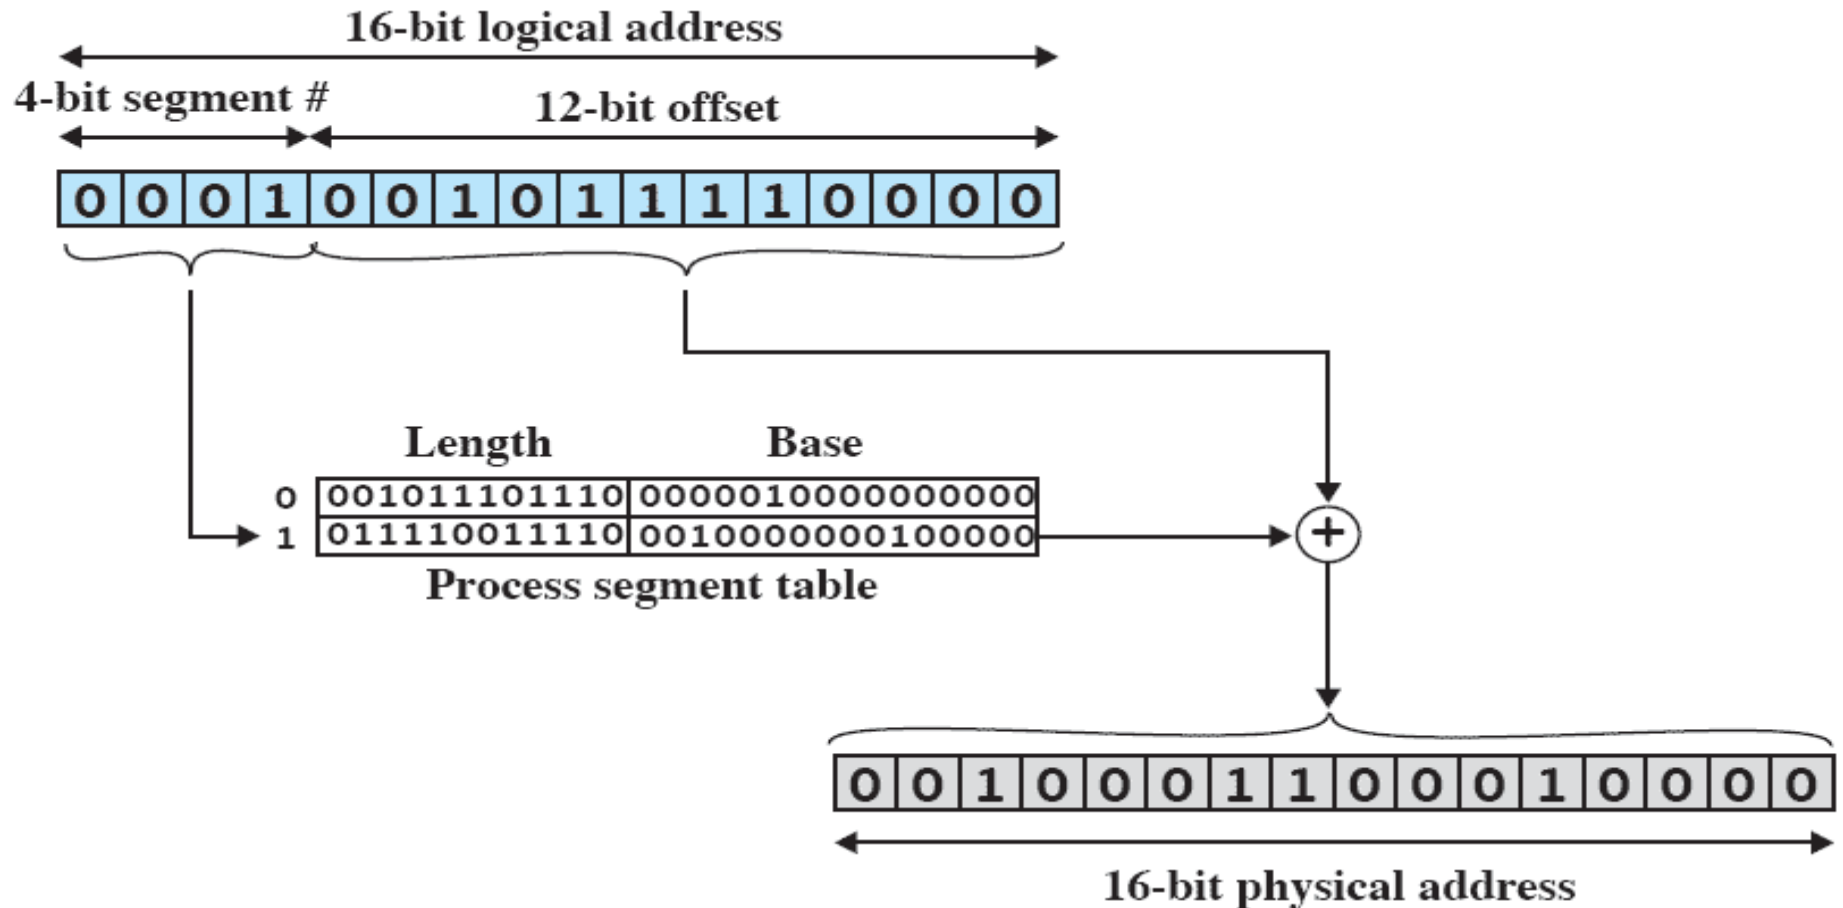
# Paging



(a) Paging



# Segmentation



(b) Segmentation

Figure 7.12 Examples of Logical-to-Physical Address Translation