

Virtual Memory

Lecture 16

Cache

- With fully associative and set associative cache, a *replacement policy* is invoked when it becomes necessary to evict a block from cache.
 - An optimal *replacement policy* would be able to look into the future to see *which blocks won't be needed* for the longest period of time.
- *Least recently used* (LRU) algorithm
 - keeps track of the last time that a block was assessed and evicts the block that has been unused for the longest period of time.
- The disadvantage of this approach is its complexity:
 - LRU has to maintain an access history for each block, which ultimately slows down the cache.

Replacement Policies

- *First-in, first-out* (FIFO) is a popular cache replacement policy.
 - The block that has been in the cache the longest, regardless of when it was last used.
- A *random* replacement policy does what its name implies:
 - It picks a block at random and replaces it with a new block.
 - Random replacement can certainly evict a block that will be needed often or needed soon, but it never.

Cache write Policies

- Cache replacement policies must also take into account *dirty blocks*,
 - Those blocks that have been updated while they were in the cache.
 - Dirty blocks must be written back to memory
 - A *write policy* determines how this will be done.
- There are two types of write policies,
 - *write through* and *write back*.
- Write through:
 - updates cache and main memory simultaneously on every write.

Cache write Policies

- Write back (also called *copyback*):
 - Updates memory only when the block is selected for replacement.
- The disadvantage of write through:
 - Memory must be updated with each cache write, which slows down the access time on updates.
 - This slowdown is usually negligible, because the majority of accesses tend to be reads, not writes.
- The advantage of write back:
 - Memory traffic is minimized,
- Its disadvantage:
 - Memory does not always agree with the value in cache, causing problems in systems with many concurrent users.

Virtual Memory

- Cache memory:
 - Enhances performance by providing faster memory access speed.
- Virtual memory:
 - Enhances performance by providing greater memory capacity, without the expense of adding main memory.
- Instead, a portion of a disk drive serves as an extension of main memory.

Virtual Memory

- The amount of real memory in a computer is limited to the amount of RAM installed.
 - Common memory sizes are 256MB, 512MB, and 1GB.
- Since it has a finite amount of RAM:
 - It is possible to run out of memory when too many programs are running at one time
 - if you load the operating system, an e-mail program, a Web browser and word processor into RAM simultaneously, RAM may not be enough to hold it all
- This is where virtual memory comes in:
 - Virtual memory increases the available memory of computer by enlarging the "address space," or places in memory where data can be stored

Virtual Memory Terminologies

- **Paging:**
 - One of the memory-management schemes by which a computer can store and retrieve data from hard disk for use in main memory
- The operating system retrieves data from hard disk in same-size blocks called *pages*
 - The advantage of paging:
 - It allows the physical address space of a process to be noncontiguous
- **Page Fault**
 - when a program tries to access pages that are not currently mapped to physical memory. This is known as a page fault.

Virtual Memory Terminologies


- Page Frame:
 - If a system uses paging, **virtual memory** partitions main memory into individually managed *page frames*:
 - That are written (*or paged*) to disk when they are not immediately needed.
- A *physical address*:
 - The actual memory address of physical memory.

Virtual Memory- Terminologies

- **Programs** create **virtual addresses**:
 - That are mapped to physical addresses by the memory manager.
- Page Table
 - Information concerning the location of each page, whether on disk or in memory, is maintained in a data structure called a page table.
 - There is one page table for each active process.

Page Table

- Each process has its own page table
 - Each page table entry contains the frame number of the corresponding page in main memory
 - A bit is needed to indicate whether the page is in main memory or not



(a) Paging only


Virtual Memory

- A *physical address* is the actual memory address of physical memory.
 - Each program has its own virtual space, which is the set of addresses that a program uses for load and store operation
- The physical address space is the set of addresses used to reference location in the main memory
 - Programs create *virtual addresses* that are *mapped* to physical addresses by the memory manager.

Virtual Memory

- Virtual and physical address are used to describe address in virtual and physical address space
 - Virtual address space is divided in pages – hard disk
 - Virtual page
 - Some of the pages are copied into memory – **page frame**
 - Physical page

Virtual memory



Virtual Memory Organization


- Main memory and virtual memory are **divided** into **equal sized pages**.
- The entire address space required by a process need not be in memory at once.
 - Some parts can be on disk, while others are in main memory.
 - The pages allocated to a process do not need to be stored contiguously-- either on disk or in memory.
- Only the needed pages are in memory at any time,
 - the unnecessary pages are in slower disk storage.

Address Translation

- When a process generates a **virtual address**, the operating system translates it into a **physical memory address**.
- To accomplish this,
 - The virtual address is divided into two fields:
 - A **page field**, and an **offset field**.
 - The page field determines the page location of the address, and
 - The offset indicates the location of the address within the page.
- The logical page number is translated into a physical page frame through a lookup in the page table.

Page Table


- The entire page table may take up too much main memory
 - Page tables are also stored in virtual memory
- When a process is running, part of its page table is in main memory



COA, MESE 15B, Asst Prof Athar Mohsin, MCS-NUST

17

Page Tables



COA, MESE 15B, Asst Prof Athar Mohsin, MCS-NUST


18

Page Fault

- If the valid bit is zero in the page table entry for the logical address, this means that the page is not in memory and must be fetched from disk.
 - This is a page fault.
 - If necessary, a page is evicted from memory and is replaced by the page retrieved from disk, and the valid bit is set to 1.
- If the valid bit is 1, the virtual page number is replaced by the physical frame number.
- The data is then accessed by adding the offset to the physical frame number.

Virtual Memory-Example

- suppose a system has a virtual address space of **8K** and a physical address space of **4K**, and the system uses byte addressing and the page size is 1024.
 - So $2^{13}/2^{10} = 2^3$ virtual pages.
- A virtual address has 13 bits ($8K = 2^{13}$) with 3 bits for the page field and 10 for the offset,
- A physical memory address requires 12 bits, the first two bits for the page frame and the trailing 10 bits the offset.



Virtual Memory-Example

- If we have the page table shown below.
 - What happens when CPU generates address $5459_{10} = 1010101010011_2$?

	Page	Frame	Valid Bit	Addresses
Page Table	0	-	0	Page 0 : 0 - 1023
	1	3	1	1 : 1024 - 2047
	2	0	1	2 : 2048 - 3071
	3	-	0	3 : 3072 - 4095
	4	-	0	4 : 4096 - 5119
	5	1	1	5 : 5120 - 6143
	6	2	1	6 : 6144 - 7167
	7	-	0	7 : 7168 - 8191

COA, MESE 15B, Asst Prof Athar Mohsin, MCS-NUST

21

Virtual Memory

- The address 1010101010011_2 is converted to physical address 010101010011 because the page field 101 is replaced by frame number 01 through a lookup in the page table.

	Page	Frame	Valid Bit	Addresses
Page Table	0	-	0	Page 0 : 0 - 1023
	1	3	1	1 : 1024 - 2047
	2	0	1	2 : 2048 - 3071
	3	-	0	3 : 3072 - 4095
	4	-	0	4 : 4096 - 5119
	5	1	1	5 : 5120 - 6143
	6	2	1	6 : 6144 - 7167
	7	-	0	7 : 7168 - 8191

COA, MESE 15B, Asst Prof Athar Mohsin, MCS-NUST

22

Virtual Memory- Example

- What happens when the CPU generates address 1000000000100_2 ?


		Frame	Valid Bit	Addresses	
Page Table	Page 0	-	0	Page 0 :	0 - 1023
	1	3	1	1 :	1024 - 2047
	2	0	1	2 :	2048 - 3071
	3	-	0	3 :	3072 - 4095
	4	-	0	4 :	4096 - 5119
	5	1	1	5 :	5120 - 6143
	6	2	1	6 :	6144 - 7167
	7	-	0	7 :	7168 - 8191

Page Faults (like “Cache Misses”)


- What if an object is on disk rather than in memory?
 - Page table entry indicates virtual address not in memory
 - OS exception handler invoked to move data from disk into memory

- current process suspends, others can resume
- OS has full control over placement, etc.

Before fault



After fault



Virtual Memory

- virtual memory is also a factor in the calculation, to consider page table access time. Example
- Suppose a main memory access takes 200ns, the page fault rate is 1%, and it takes 10ms to load a page from disk. We have:

$$\text{EAT} = 0.99(200\text{ns} + 200\text{ns}) + 0.01(10\text{ms}) = 100,396\text{ns}.$$

- Even if there had no page faults, the EAT would be 400ns because memory is always read twice:
 - First to access the page table, and second to load the page from memory.