### CRASH RECOVERY

Instructor

Dr. Sanam Shahla Rizvi
PhD in Information and Communication
from Ajou University, Korea

#### NEED FOR RECOVERY

- Natural physical disasters
  - Fires, floods, earthquakes or power outage
- Sabotage
  - Intentional contamination or destruction of data, hardware, software facilities
- Carelessness
  - Unintentional contamination or destruction of data, hardware, software facilities
- Disk malfunctions
  - Head crashes, defective disks, unreadable tracks
- System crashes
  - Hardware malfunctions, resulting in loss of main and cache memory
- System software errors
  - Abnormal termination or destruction of DBMS
- Application software errors
  - Logical errors in program accessing database

#### RECOVERY MANAGER

- When a DMBS restarts after crashes, the recovery manager given control to must bring the database on a consistent state.
- Ensures atomicity by undoing the actions that do not commit.
- Ensures durability by making sure that all actions of committed transactions survive system crashes.

# RECOVERY RELATED STEPS DURING NORMAL EXECUTION

- Log of all modifications saves on stable storage.
- Maintain multiple copies on different locations on nonvolatile storage devices.
- Write log entries in stable storage before the transaction commits.

#### ARIES

- Analysis: Identifies dirty pages in the buffer pool (i.e., changes that have not been written to disk) and active transactions at the time of the crash.
- **Redo:** Repeats all actions, starting from an appropriate point in the log, and restores the database state to what it was at the time of the crash.
- Undo: Undoes the actions of transactions that did not commit, so that the database reflects only the actions of committed transactions.

#### ARIES PRINCIPALS

- 1. Write-ahead logging (WAL): Any change to a database object is first recorded in the log; the record in the log must be written to stable storage before the change to the database object is written to disk.
- 2. Repeating history during Redo: Upon restart following a crash, ARIES retraces all actions of the DBMS before the crash and brings the system back to the exact state that it was in at the time of the crash. Then, it undoes the actions of transactions that were still active at the time of the crash (effectively aborting them).
- 3. Logging changes during Undo: Changes made to the database while undoing a transaction are logged in order to ensure that such an action is not repeated in the event of repeated (failures causing) restarts.

## Log/Trail/Journal

- Log record is written for each of following actions:
  - Updating a page
  - Commit
  - Abort
  - End
  - Undoing an update

#### RECOVERING FROM A SYSTEM CRASH

#### o Analysis:

- It determines the point in the log at which to start the Redo pass.
- It determines (a conservative superset of the) pages in the buffer pool that were dirty at the time of the crash.
- It identifies transactions that were active at the time of the crash and must be undone.

#### • Redo:

- The logged actions are reapplied.
- The page is set redone log record.

#### Undo

- Scan backward from end of log.
- Undo actions of transactions active lastly.

```
• LSN
              LOG
              begin
o 00
              update: T1 writes P5
10
              update: T2 writes P3
o 20
              T1 aborts
o 30
              CLR: Undo T1 LSN10, T1 end
o 40
              update:T3 writes P1
o 50
o 60
              update: T2 writes P5
       CRASH, RESTART
o 70
o 80
       CRASH, RESTART
o 90
```

**o** 90

• LSN	<u>LOG</u>
<b>o</b> 00	begin
<b>o</b> 10	update: T1 writes P5
<b>o</b> 20	update: T2 writes P3
<b>o</b> 30	T1 aborts
<b>o</b> 40	CLR: Undo T1 LSN10, T1 end
<b>o</b> 50	update:T3 writes P1
<b>o</b> 60	update: T2 writes P5
0	CRASH, RESTART
<b>o</b> 70	CLR: Undo T2 LSN60
<b>o</b> 80	CLR: Undo T3 LSN50, T3 end
0	CRASH, RESTART

CLR: Undo T2 LSN20, T2 end



